

The Transformer Architecture Breakdown: A Mathematical Walkthrough

-Pradeep Rajasekar(Ajish)

Contents

Introduction	2
Understanding Transformers	2
What is a Transformer?	2
Why Transformers?	2
Part I: The Encoder	3
Step 1: Tokenization and Input Representation	3
Vocabulary and One-Hot Encoding	3
Embedding Layer	3
Step 2: Positional Encoding	4
Positional Encoding Calculation	4
Combining Embeddings and Positional Encoding	4
Step 3: Multi-Head Self-Attention	5
Linear Projections to Query, Key, and Value	5
Scaled Dot-Product Attention	5
Multi-Head Attention	6
Step 4: Add & Norm	7
Residual Connection	7
Layer Normalization	7
Step 5: Position-Wise Feed-Forward Network	7
Feed-Forward Network (FFN)	7
Add & Norm	8
Step 6: Stacking Multiple Encoder Layers	8
Part II: The Decoder	8
Step 7: Shifted Right Input and Embedding	8
Preparing the Decoder Input	8
Embedding and Positional Encoding	9
Step 8: Masked Multi-Head Self-Attention	9
Masking Mechanism	9
Applying the Mask	9
Step 9: Encoder-Decoder Attention	10
Computing Queries, Keys, and Values	10
Attention Calculation	10
Step 10: Add & Norm and Feed-Forward Network	11
Add & Norm (Encoder-Decoder Attention)	11
Position-Wise Feed-Forward Network	11
Step 11: Stacking Multiple Decoder Layers	11
Step 12: Final Linear Layer and Softmax	11
Linear Projection	12
Softmax Function	12
Conclusion	12

Introduction

The advent of the Transformer architecture has marked a significant milestone in the field of natural language processing (NLP). Introduced by Vaswani et al. in 2017, Transformers have revolutionized how models understand and generate language by relying entirely on attention mechanisms, eschewing the recurrence and convolution used in previous architectures like recurrent neural networks (RNNs) and convolutional neural networks (CNNs). This shift has enabled models to capture long-range dependencies more effectively and parallelize computations, leading to substantial improvements in performance and efficiency.

In this comprehensive guide, we delve deep into the mathematical foundations of the Transformer model. We will walk through each component, from the encoder's input embeddings to the decoder's output generation, using a simple yet illustrative example:

- **Input Sentence:** "Ajish works as an AI"
- **Predicted Next Word:** "Engineer"

Our aim is to not only dissect the mathematical operations involved in each step but also to explain the rationale behind them and how they contribute to the overall architecture's efficacy. By the end of this walkthrough, you will have a thorough understanding of how each part of the Transformer works and why it is essential for the model's ability to process and generate language.

Understanding Transformers

Before we dive into the detailed walkthrough, let's establish a foundational understanding of what Transformers are and why they have become pivotal in NLP.

What is a Transformer?

A Transformer is a type of neural network architecture that relies entirely on self-attention mechanisms to process sequences of data. Unlike traditional sequence models like RNNs and CNNs, Transformers do not require sequential data to be processed in order. Instead, they process all tokens in the input sequence simultaneously, allowing for greater parallelization and efficiency during training.

The Transformer architecture is composed of an encoder and a decoder, both of which are stacks of identical layers. The encoder processes the input sequence and generates a rich representation, while the decoder uses this representation to generate the output sequence, one token at a time.

Why Transformers?

Traditional models like RNNs suffer from limitations in capturing long-range dependencies due to vanishing gradients and the inherent sequential nature of their computations, which makes parallelization difficult. Transformers address these issues by:

- **Utilizing Self-Attention:** Self-attention mechanisms allow the model to weigh the importance of different words in a sequence relative to each other, regardless of their position. This enables the model to capture long-range dependencies more effectively.
- **Parallel Processing:** By processing all tokens simultaneously, Transformers enable faster training times and more efficient use of computational resources.
- **Capturing Global Dependencies:** The self-attention mechanism provides a way for the model to consider the entire sequence context for each token, improving the model's understanding of the language structure.

These advantages have made Transformers the foundation for many state-of-the-art models in NLP, such as BERT and the GPT series.

Part I: The Encoder

The encoder's primary function is to transform the input sequence into a rich, contextualized representation that encapsulates the meaning and relationships between the tokens. Let's explore each step in detail.

Step 1: Tokenization and Input Representation

Vocabulary and One-Hot Encoding

Tokenization is the process of breaking down a sentence into individual words or tokens. For our example:

- **Input Sentence:** "Ajish works as an AI"

After tokenization, we have the tokens:

$$\text{Tokens} = [\text{"Ajish"}, \text{"works"}, \text{"as"}, \text{"an"}, \text{"AI"}]$$

Next, we need to map these tokens to numerical representations that the model can process.

Vocabulary (V): The vocabulary is a collection of all unique tokens that the model recognizes. For simplicity, we'll use a small vocabulary of 10 words:

$$V = \{\text{"Ajish"}, \text{"works"}, \text{"as"}, \text{"an"}, \text{"AI"}, \text{"Engineer"}, \text{"the"}, \text{"is"}, \text{"a"}, \text{"."}\}$$

One-Hot Encoding: Each word in the vocabulary is represented as a one-hot vector, a binary vector of size $|V|$ (the size of the vocabulary), where only the index corresponding to the word is 1, and all others are 0.

For example:

- **"Ajish":** $\mathbf{o}_{\text{Ajish}} = [1, 0, 0, 0, 0, 0, 0, 0, 0, 0]$
- **"AI":** $\mathbf{o}_{\text{AI}} = [0, 0, 0, 0, 1, 0, 0, 0, 0, 0]$

Why One-Hot Encoding?

One-hot encoding provides a simple way to represent categorical data numerically. However, it has limitations:

- **Sparsity:** The vectors are high-dimensional and sparse, which is computationally inefficient.
- **Lack of Semantic Information:** One-hot vectors do not capture any semantic relationships between words. For instance, "AI" and "Engineer" are related concepts but their one-hot vectors are orthogonal.

Embedding Layer

To address the limitations of one-hot encoding, we use an **embedding layer** to project the sparse one-hot vectors into a dense, continuous vector space where semantic relationships can be captured.

Embedding Matrix (\mathbf{W}_{emb}): A learnable weight matrix of dimensions $|V| \times d_{\text{model}}$, where d_{model} is the model's hidden size (e.g., 512). Each row corresponds to a word's embedding vector.

Computing Word Embeddings:

For a word i (with one-hot vector \mathbf{o}_i):

$$\mathbf{e}_i = \mathbf{o}_i \mathbf{W}_{\text{emb}}$$

Since \mathbf{o}_i is a one-hot vector, this operation effectively selects the i -th row of \mathbf{W}_{emb} .

Purpose of Embeddings:

- **Semantic Representation:** Embeddings allow words with similar meanings to have similar representations. For example, "AI" and "Engineer" might have embeddings that are close in the vector space.

- **Dimensionality Reduction:** Embeddings reduce the high-dimensional one-hot vectors to lower-dimensional dense vectors, making computations more efficient.

For our example, the embeddings for the words in the input sentence are:

$$\begin{aligned} \mathbf{e}_{\text{Ajish}} &= \mathbf{o}_{\text{Ajish}} \mathbf{W}_{\text{emb}} \\ \mathbf{e}_{\text{works}} &= \mathbf{o}_{\text{works}} \mathbf{W}_{\text{emb}} \\ \mathbf{e}_{\text{as}} &= \mathbf{o}_{\text{as}} \mathbf{W}_{\text{emb}} \\ \mathbf{e}_{\text{an}} &= \mathbf{o}_{\text{an}} \mathbf{W}_{\text{emb}} \\ \mathbf{e}_{\text{AI}} &= \mathbf{o}_{\text{AI}} \mathbf{W}_{\text{emb}} \end{aligned}$$

These embeddings capture the semantic meaning of each word and are the starting point for the model's processing.

Step 2: Positional Encoding

Transformers process all tokens simultaneously and, therefore, lack inherent positional information about the order of the tokens in the sequence. To inject this information, we use **positional encoding**.

Positional Encoding Calculation

The positional encoding adds unique positional information to each token embedding, enabling the model to understand the order of the tokens.

Formula:

For position pos (starting from 0) and embedding dimension i (ranging from 0 to $d_{\text{model}} - 1$):

$$\begin{aligned} \text{PE}_{(pos, 2i)} &= \sin\left(\frac{pos}{10000^{2i/d_{\text{model}}}}\right), \\ \text{PE}_{(pos, 2i+1)} &= \cos\left(\frac{pos}{10000^{2i/d_{\text{model}}}}\right). \end{aligned}$$

Explanation:

- **Sine and Cosine Functions:** Using sine and cosine functions with varying frequencies allows each position to have a unique encoding. The even dimensions use sine functions, and the odd dimensions use cosine functions.
- **Frequency Variation:** The term $\frac{1}{10000^{2i/d_{\text{model}}}}$ scales the frequency based on the dimension i . This creates a spectrum of frequencies across the dimensions, allowing the model to learn position-related patterns.

Combining Embeddings and Positional Encoding

The final input representation for each token is the sum of its word embedding and positional encoding:

$$\mathbf{x}_i = \mathbf{e}_i + \mathbf{p}_i,$$

where \mathbf{p}_i is the positional encoding vector for position i .

Why Add Positional Encoding?

- **Sequence Order Information:** Adding positional encodings enables the model to distinguish between tokens at different positions.
- **Smoothness:** The continuous nature of sine and cosine functions ensures that positions close to each other have similar encodings, which is beneficial for learning.

Example:

For the word “Ajish” at position 0:

- **Embedding:** $\mathbf{e}_{\text{Ajish}}$
- **Positional Encoding:** \mathbf{p}_0
- **Input Representation:** $\mathbf{x}_0 = \mathbf{e}_{\text{Ajish}} + \mathbf{p}_0$

This process is repeated for each token in the input sequence.

Step 3: Multi-Head Self-Attention

Self-attention allows the model to weigh the influence of other tokens in the sequence when encoding a particular token, capturing dependencies regardless of their distance apart. The multi-head mechanism enhances this capability by allowing the model to focus on different aspects of the token relationships simultaneously.

Linear Projections to Query, Key, and Value

For each token, we compute three vectors: **Query** (\mathbf{q}_i), **Key** (\mathbf{k}_i), and **Value** (\mathbf{v}_i).

Linear Transformations:

For each token representation \mathbf{x}_i :

$$\begin{aligned}\mathbf{q}_i &= \mathbf{x}_i \mathbf{W}_Q, \\ \mathbf{k}_i &= \mathbf{x}_i \mathbf{W}_K, \\ \mathbf{v}_i &= \mathbf{x}_i \mathbf{W}_V,\end{aligned}$$

where:

- $\mathbf{W}_Q, \mathbf{W}_K, \mathbf{W}_V$ are learnable weight matrices of dimensions $d_{\text{model}} \times d_k$ (e.g., 512×64 if $d_k = 64$).

Purpose:

- **Query** (\mathbf{q}_i): Represents the current token we are focusing on.
- **Key** (\mathbf{k}_i): Represents the tokens we compare against.
- **Value** (\mathbf{v}_i): Contains the information to be aggregated based on attention scores.

Scaled Dot-Product Attention

The attention mechanism computes a weighted sum of the values, where the weights are determined by the compatibility of the queries and keys.

Attention Scores Calculation:

For all tokens, we compute the attention scores matrix \mathbf{S} :

$$\mathbf{S} = \mathbf{Q}\mathbf{K}^\top,$$

where:

- \mathbf{Q} is the matrix of queries ($n \times d_k$).
- \mathbf{K}^\top is the transpose of the keys matrix ($d_k \times n$).
- The result is an $n \times n$ matrix of scores, where S_{ij} represents the score between token i and token j .

Scaling Factor:

To prevent the dot product values from becoming too large, which can push the softmax into regions with very small gradients, we scale the scores:

$$\mathbf{S}_{\text{scaled}} = \frac{\mathbf{S}}{\sqrt{d_k}}.$$

Softmax Normalization:

We apply the softmax function to each row to obtain the attention weights:

$$\mathbf{A}_{ij} = \text{softmax}(\mathbf{S}_{\text{scaled}})_{ij} = \frac{\exp(S_{\text{scaled},ij})}{\sum_{k=1}^n \exp(S_{\text{scaled},ik})}.$$

Attention Output:

The output for each token is a weighted sum of the value vectors:

$$\mathbf{z}_i = \sum_{j=1}^n A_{ij} \mathbf{v}_j.$$

Interpretation:

- The attention weights A_{ij} indicate how much attention the model pays to token j when processing token i .
- By summing over all tokens, the model integrates information from the entire sequence, weighted by their relevance.

Multi-Head Attention

Instead of performing a single attention function, the Transformer uses multiple attention heads to capture different types of relationships.

Splitting into Multiple Heads:

- The model splits the d_{model} dimensional space into h smaller subspaces (e.g., $h = 8$, so each head has dimension $d_k = d_{\text{model}}/h = 64$).

Per-Head Computation:

For each head i :

$$\text{head}_i = \text{Attention}(\mathbf{Q}_i, \mathbf{K}_i, \mathbf{V}_i),$$

where $\mathbf{Q}_i, \mathbf{K}_i, \mathbf{V}_i$ are the queries, keys, and values for head i .

Concatenation and Final Linear Transformation:

The outputs of the heads are concatenated and projected back to d_{model} dimensions:

$$\mathbf{Z} = \text{Concat}(\text{head}_1, \dots, \text{head}_h) \mathbf{W}_O,$$

where \mathbf{W}_O is a learnable weight matrix of size $(h \cdot d_k) \times d_{\text{model}}$.

Purpose of Multi-Head Attention:

- **Multiple Perspectives:** Each head can focus on different types of relationships (e.g., syntactic vs. semantic).
- **Enhanced Representational Power:** By combining information from multiple heads, the model can capture complex patterns more effectively.

Example:

In our example sentence, one head might focus on the subject-verb relationship between “Ajish” and “works,” while another head might focus on the relationship between “an” and “AI.”

Step 4: Add & Norm

After the multi-head attention, we apply a residual connection followed by layer normalization to stabilize and improve the training of the network.

Residual Connection

We add the original input \mathbf{X} to the output of the multi-head attention \mathbf{Z} :

$$\mathbf{X}' = \mathbf{X} + \mathbf{Z}.$$

Purpose of Residual Connections:

- **Mitigating Vanishing Gradients:** They help gradients flow through the network more effectively, especially in deep networks.
- **Ease of Optimization:** Residual connections allow the model to learn incremental changes over the identity mapping, simplifying the learning process.

Layer Normalization

We apply layer normalization to \mathbf{X}' :

$$\mathbf{X}'' = \text{LayerNorm}(\mathbf{X}'),$$

where layer normalization normalizes the inputs across the features (dimensions) for each token.

Purpose of Layer Normalization:

- **Stabilizing Learning:** It reduces internal covariate shift by ensuring that the input to each layer maintains a consistent distribution.
- **Faster Convergence:** Models with layer normalization tend to converge faster during training.

Step 5: Position-Wise Feed-Forward Network

Each token's representation is further transformed through a fully connected feed-forward network applied identically to each position.

Feed-Forward Network (FFN)

The FFN consists of two linear transformations with a ReLU activation in between:

$$\mathbf{F}(\mathbf{x}) = \text{ReLU}(\mathbf{x}\mathbf{W}_1 + \mathbf{b}_1)\mathbf{W}_2 + \mathbf{b}_2.$$

Parameters:

- \mathbf{W}_1 and \mathbf{W}_2 are weight matrices of dimensions $d_{\text{model}} \times d_{\text{ff}}$ and $d_{\text{ff}} \times d_{\text{model}}$, respectively.
- \mathbf{b}_1 and \mathbf{b}_2 are bias vectors.
- d_{ff} is typically larger than d_{model} (e.g., 2048), allowing the network to learn more complex transformations.

Activation Function:

- **ReLU** introduces non-linearity, enabling the network to learn complex patterns and interactions between features.

Purpose of FFN:

- **Complex Feature Transformation:** The FFN allows the model to apply position-wise transformations to capture complex relationships that may not be captured by the attention mechanism alone.
- **Depth of Processing:** By applying non-linear transformations, the model can learn higher-order features.

Add & Norm

After the FFN, we apply another residual connection and layer normalization:

$$\mathbf{Y} = \text{LayerNorm}(\mathbf{X}'' + \mathbf{F}(\mathbf{X}'')).$$

Purpose:

- **Residual Connection:** As before, it helps with gradient flow and learning incremental transformations.
- **Layer Normalization:** Ensures stable distributions of activations, aiding in training.

Step 6: Stacking Multiple Encoder Layers

The above processes (Steps 3 to 5) are encapsulated in an **encoder layer**. The Transformer stacks multiple such layers (N , e.g., 6 layers) to build a deep encoder.

Purpose of Stacking Layers:

- **Hierarchical Feature Learning:** Each layer can build upon the representations learned by the previous layer, capturing increasingly abstract features.
- **Deep Processing:** Multiple layers allow the model to perform complex transformations and capture intricate patterns in the data.

Sequential Processing:

- The output of one layer serves as the input to the next.
- This process allows the model to refine token representations progressively.

Part II: The Decoder

The decoder generates the output sequence by attending to both the encoder's output representations and previously generated tokens. Let's explore how the decoder functions in detail.

Step 7: Shifted Right Input and Embedding*Preparing the Decoder Input*

During Training:

- **Target Sequence:** "Ajish works as an AI Engineer"
- **Shifted Input:** We shift the target sequence to the right by one position to create the decoder input.

Original Target = ["Ajish", "works", "as", "an", "AI", "Engineer"]

Shifted Input = [" $\text{\textcircled{S}}$ OS", "Ajish", "works", "as", "an", "AI"]

- **$\text{\textcircled{S}}$ OS (Start-of-Sequence):** A special token indicating the beginning of the sequence.

Purpose of Shifting:

- **Teacher Forcing:** During training, we provide the correct previous tokens to the decoder to learn the correct next token.
- **Autoregressive Modeling:** The decoder predicts the next token based on the previous tokens.

During Inference:

- The decoder generates one token at a time, using its own previous predictions as input.

Embedding and Positional Encoding

Each token in the decoder input is embedded and combined with positional encoding, similar to the encoder steps.

Embedding:

For each token i :

$$\mathbf{e}_{\text{dec},i} = \mathbf{o}_{\text{dec},i} \mathbf{W}_{\text{emb}},$$

where $\mathbf{o}_{\text{dec},i}$ is the one-hot vector for the decoder input token, and \mathbf{W}_{emb} is the same embedding matrix used in the encoder.

Positional Encoding:

We compute positional encodings \mathbf{p}_i for the decoder positions, as before.

Combining Embeddings and Positional Encoding:

$$\mathbf{x}_{\text{dec},i} = \mathbf{e}_{\text{dec},i} + \mathbf{p}_i.$$

Purpose:

- **Consistency:** Using the same embedding matrix ensures that the embeddings are in the same vector space.
- **Order Information:** Positional encodings allow the decoder to understand the order of the generated tokens.

Step 8: Masked Multi-Head Self-Attention

The decoder's self-attention mechanism is modified to prevent it from accessing future tokens, ensuring the model generates outputs causally.

Masking Mechanism

Causal Mask (M):

The mask is an $n \times n$ matrix (for a sequence of length n) defined as:

$$M_{ij} = \begin{cases} 0 & \text{if } j \leq i, \\ -\infty & \text{if } j > i. \end{cases}$$

Application:

When computing attention scores, we add the mask to the scaled dot-product attention scores:

$$\mathbf{S}_{\text{masked}} = \frac{\mathbf{Q}_{\text{dec}} \mathbf{K}_{\text{dec}}^\top}{\sqrt{d_k}} + \mathbf{M}.$$

Applying the Mask

Softmax with Mask:

The $-\infty$ values in the mask ensure that after the softmax function, the attention weights for future tokens are zero.

$$\mathbf{A}_{\text{masked},ij} = \frac{\exp(S_{\text{masked},ij})}{\sum_{k=1}^i \exp(S_{\text{masked},ik})}.$$

Attention Output:

$$\mathbf{z}_{\text{dec},i} = \sum_{j=1}^i \mathbf{A}_{\text{masked},ij} \mathbf{v}_{\text{dec},j}.$$

Purpose of Masking:

- **Autoregressive Property:** Ensures that the model cannot “cheat” by looking at future tokens during training.
- **Sequential Generation:** Mimics the conditions during inference, where future tokens are not available.

Step 9: Encoder-Decoder Attention

This layer allows the decoder to incorporate information from the encoder’s output, effectively focusing on relevant parts of the input sequence when generating each output token.

Computing Queries, Keys, and Values

Queries from Decoder Output:

$$\mathbf{Q}_{\text{dec}} = \mathbf{Y}_{\text{dec}} \mathbf{W}_Q^{\text{dec}}.$$

Keys and Values from Encoder Output:

$$\mathbf{K}_{\text{enc}} = \mathbf{Y}_{\text{enc}} \mathbf{W}_K^{\text{enc}}, \quad \mathbf{V}_{\text{enc}} = \mathbf{Y}_{\text{enc}} \mathbf{W}_V^{\text{enc}}.$$

Dimensions:

- \mathbf{Y}_{dec} : Output from the decoder’s masked self-attention and Add & Norm.
- \mathbf{Y}_{enc} : Final output from the encoder.

Attention Calculation

Attention Scores:

$$\mathbf{S}_{\text{enc-dec}} = \frac{\mathbf{Q}_{\text{dec}} \mathbf{K}_{\text{enc}}^{\top}}{\sqrt{d_k}}.$$

Softmax Normalization:

$$\mathbf{A}_{\text{enc-dec}} = \text{softmax}(\mathbf{S}_{\text{enc-dec}}).$$

Attention Output:

$$\mathbf{Z}_{\text{enc-dec}} = \mathbf{A}_{\text{enc-dec}} \mathbf{V}_{\text{enc}}.$$

Purpose:

- **Incorporating Encoder Context:** Allows the decoder to attend to specific parts of the input sequence, providing context for generating the next token.
- **Aligning Input and Output:** Helps the model learn the alignment between input and output sequences.

Example:

When predicting “Engineer,” the decoder can attend to “AI” in the encoder’s output, recognizing that “AI” is related to “Engineer.”

Step 10: Add & Norm and Feed-Forward Network

Add & Norm (Encoder-Decoder Attention)

Residual Connection:

$$\mathbf{X}'' = \mathbf{Y}_{\text{dec}} + \mathbf{Z}_{\text{enc-dec}}.$$

Layer Normalization:

$$\mathbf{X}''' = \text{LayerNorm}(\mathbf{X}'').$$

Purpose:

- **Combining Information:** Merges the decoder's self-attention output with the encoder-decoder attention output.
- **Stabilization:** Layer normalization ensures stable gradients and consistent distributions.

Position-Wise Feed-Forward Network

Feed-Forward Transformation:

$$\mathbf{F}_{\text{dec}}(\mathbf{x}) = \text{ReLU}(\mathbf{x}\mathbf{W}_1^{\text{dec}} + \mathbf{b}_1^{\text{dec}})\mathbf{W}_2^{\text{dec}} + \mathbf{b}_2^{\text{dec}}.$$

Add & Norm:

$$\mathbf{Y}'_{\text{dec}} = \text{LayerNorm}(\mathbf{X}''' + \mathbf{F}_{\text{dec}}(\mathbf{X}''')).$$

Purpose:

- **Refinement:** The FFN allows the decoder to apply further transformations to the combined representation.
- **Non-Linearity:** The ReLU activation introduces non-linearity, enabling complex pattern learning.

Step 11: Stacking Multiple Decoder Layers

Similar to the encoder, the decoder layers are stacked N times to enhance the model's capacity to capture complex patterns and dependencies.

Purpose:

- **Deep Understanding:** Multiple layers allow the decoder to perform deep transformations, refining its outputs.
- **Iterative Refinement:** Each layer builds upon the previous, improving the quality of the generated tokens.

Sequential Processing:

- The output of one decoder layer becomes the input to the next.
- This allows for multiple rounds of attention over the encoder's output.

Step 12: Final Linear Layer and Softmax

The decoder's final output is transformed into probabilities over the vocabulary to generate the next word.

Linear Projection

Transformation:

$$\mathbf{L} = \mathbf{Y}'_{\text{dec}} \mathbf{W}_{\text{out}} + \mathbf{b}_{\text{out}},$$

where:

- \mathbf{W}_{out} is a weight matrix of size $d_{\text{model}} \times |V|$.
- \mathbf{b}_{out} is a bias vector of size $|V|$.
- The output \mathbf{L} is a matrix of size $n \times |V|$, where each row corresponds to a token's unnormalized log probabilities over the vocabulary.

Softmax Function

Computing Probabilities:

For each token position i :

$$P_i = \text{softmax}(\mathbf{L}_i),$$

where \mathbf{L}_i is the i -th row of \mathbf{L} .

Selecting the Next Word:

- The word with the highest probability at position i is selected as the predicted next token.
- In our example, for the position following “AI,” the model predicts “**Engineer**”.

Purpose:

- **Probability Distribution:** Converts the decoder's output into a probability distribution over the vocabulary.
- **Word Prediction:** Enables the selection of the most probable next word.

Example:

Suppose the softmax outputs for the vocabulary are:

- “**Engineer**”: 0.6
- “**Scientist**”: 0.2
- “**Developer**”: 0.1
- **Others:** Remaining probabilities

The model predicts “Engineer” as it has the highest probability.

Conclusion

Through this comprehensive mathematical walkthrough, we have dissected the Transformer architecture's encoder and decoder components, elucidating how they collaboratively process input sequences and generate predictions. Each step—from embedding and positional encoding to multi-head attention and feed-forward networks—plays a crucial role in enabling the model to understand and generate language effectively.

• Self-Attention Mechanisms:

- Allow the model to capture dependencies between tokens, regardless of their positions in the sequence.

- Enable the model to focus on relevant parts of the sequence when processing each token.
- **Positional Encoding:**
 - Injects sequence order information, crucial for understanding language structure.
 - Ensures that tokens have unique representations based on their positions.
- **Multi-Head Attention:**
 - Enables the model to focus on different aspects of the relationships between tokens simultaneously.
 - Enhances the representational capacity of the model.
- **Residual Connections and Layer Normalization:**
 - Facilitate training deep networks by ensuring stable gradients and faster convergence.
 - Improve the flow of information through the network.
- **Stacked Layers:**
 - Increase the model's depth, allowing for more complex representations and abstractions.
 - Enable hierarchical feature learning.

The Transformer architecture's ability to capture complex dependencies and contextual relationships has made it the foundation for many state-of-the-art models in NLP, such as BERT, GPT, and others. Understanding the underlying mathematical operations not only demystifies how Transformers function but also opens avenues for further innovations and applications in natural language processing and beyond.

Practical Implications:

- **Parallelization:** Transformers can be efficiently parallelized, leading to faster training times compared to RNNs.
- **Scalability:** The architecture scales well with data and model size, contributing to the success of large models like GPT-3.
- **Versatility:** While initially designed for NLP, Transformers have been adapted for use in other domains like computer vision and reinforcement learning.

References

- Vaswani, A., et al. (2017). *Attention is All You Need*. arXiv:1706.03762.
- Devlin, J., et al. (2018). *BERT: Pre-training of Deep Bidirectional Transformers for Language Understanding*. arXiv:1810.04805.
- Radford, A., et al. (2018). *Improving Language Understanding by Generative Pre-Training*. Link.